

Convolutions and CNN Layers

Advanced Machine Learning

Sam Ogden

DRAFT

This slide deck is still under active development. Expect changes before it is finalized.

Learning objectives

- Describe how an MLP sees an image (flatten, then one weight per pixel position), and explain why that makes it a poor fit for images, and name the two properties of images (locality, repeated patterns) that convolution builds in
- Compute a convolution by hand, and say what a single filter “sees” when slid over an image
- Compute output shapes under padding and stride, and parameter counts for multi-channel layers
- Describe what pooling keeps and what it throws away
- Build the layer from scratch as an `nn.Module`, then use `nn.Conv2d` and `nn.MaxPool2d` to assemble a small CNN and trace its shapes

Images and the MLP

How an MLP actually looks at an image


A very small image

a 3 × 3 image: a vertical pipe

x_0 1	x_1 0	x_2 0
x_3 1	x_4 0	x_5 0
x_6 1	x_7 0	x_8 0

- Each pixel is 0 (dark) or 1 (bright), so we can see every number
- The label x_k in the corner of each pixel is its **position**: row by row, x_0 to x_8

Step 1: flatten



- An MLP takes a **vector**, so the grid is unrolled row by row: $[1, 0, 0, 1, 0, 0, 1, 0, 0]$
- The three pipe pixels were stacked in a column. Now they are just x_0, x_3, x_6


Step 2: a model on the flattened input

We want a model that answers one question: is there a pipe on the left? (output 1 for yes, 0 for no)

$$z = w_0x_0 + w_1x_1 + \cdots + w_8x_8 + b, \quad \text{output} = \begin{cases} 1 & z > 0 \\ 0 & \text{otherwise} \end{cases}$$


- One weight w_i **per pixel position**, plus a bias b
- Nothing here says x_0 and x_3 were ever neighbors. They are just two separate inputs

The model, drawn out




- Each line is a weight. **Bold lines** carry a pixel that is on; faint lines carry a pixel that is off
- Suppose training found: $+1$ on the pipe's pixels (x_0, x_3, x_6) , -1 everywhere else, $b = -2$
- Three bold red lines: $z = 3 - 2 = 1 > 0$, so the output is **1**

Now move the pipe one pixel



A different model for the middle




- To say “yes” here, the weights have to be different: $+1$ on x_1, x_4, x_7 instead
- This model has its own weights. It has to be trained on its own examples; **nothing it learned about the left pipe carries over**

One pipe, three positions, three models

Input	Left-pipe model	Middle-pipe model	Right-pipe model
pipe on the left	1	0	0
pipe in the middle	0	1	0
pipe on the right	0	0	1

- To spot the pipe anywhere, an MLP needs a separate set of weights per position, each learned on its own
- Why not merge them into **one** model that says “pipe anywhere”?

Why not one model for every position?




The core issue:

The MLP doesn't know what's next to what

x_0 1	x_1 0	x_2 0
x_3 1	x_4 0	x_5 0
x_6 1	x_7 0	x_8 0

It has no idea that x_0 and x_3 are vertical neighbors. If a model should use that,
we have to build it in

The same digit, moved a little



- To a person, it is obviously the same digit
- To an MLP, it is the pipe all over again: **hundreds of inputs changed**, so it has to learn the digit again at every position

Another issue: Scaling up

- A 28×28 digit is 784 inputs, so each hidden unit has **784 weights**, one per pixel position
- For a modest photo, the weights alone become a problem:

224 x 224 color photo -> 150,528 inputs

1000 hidden units -> 150,528,000 weights in the first layer alone

What do we know about images?

What do we already know about images that an MLP ignores?

Two facts about images

Locality

- A pixel's meaning depends on its **neighbors**
- Edges, corners, and strokes are made of a few nearby pixels
- A pixel far across the image tells us little about it

The same pattern shows up anywhere

- A vertical edge is a vertical edge, in the top-left or the bottom-right
- A detector that works in one place should work in every place

The Big Idea

**Look for the same small pattern
everywhere in the image, using one small
set of shared weights.**

The convolution operation

Sliding a small pattern across the image

Convolution: the full definition


For a small kernel K of size $k_h \times k_w$ slid over an input X :

$$Y[i, j] = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} K[a, b] \cdot X[i + a, j + b] + b_0$$

- K is a small grid of weights (the **kernel**, or **filter**); b_0 is a single bias
- Y is the **output**, also called a **feature map**


Breakdown: one output value, in 1D

$$Y[i] = \sum_a K[a] \cdot X[i + a]$$



- Line the kernel up with a window of X , multiply each pair, and add: $(-1) \cdot 1 + 2 \cdot 2 + (-1) \cdot 3 = 0$

Breakdown: slide the window



- The window of X under the kernel is a **patch**. In 2D it is a $k_h \times k_w$ block
- One patch gives **one number**; moving the window one step gives the next. Five inputs and a kernel of three give **three** outputs
- The kernel sums to 0, so a flat stretch of X gives 0. Only the **peak** stands out

Breakdown: what “convolution” really means here

- Strictly, this formula is **cross-correlation**. A true convolution flips the kernel first
- Deep learning libraries (and d2l) compute the formula above and call it convolution anyway
- The flip makes no practical difference: the kernel is **learned**, so a model that “wanted” the flipped one would just learn the flipped values

Worked example: our setup

input X (a vertical edge)

0	0	1	1	1
0	0	1	1	1
0	0	1	1	1
0	0	1	1	1
0	0	1	1	1

kernel K

-1	0	1
-1	0	1
-1	0	1

- X : dark on the left, bright on the right. K : *right minus left*
- Bias $b_0 = 0$ for now

Worked example: first position

input X					kernel K			patch $\times$ kernel			output Y (sum = 3)		
0	0	1	1	1	-1	0	1	0	0	1	3		
0	0	1	1	1	-1	0	1	0	0	1			
0	0	1	1	1	-1	0	1	0	0	1			
0	0	1	1	1									
0	0	1	1	1									

- The patch is the top-left 3×3 corner of X ; multiply by K and add: $0 + 0 + 1$ in each row, three rows, total 3

Worked example: slide one step right

input X					kernel K			patch $\times$ kernel			output Y (sum = 3)		
0	0	1	1	1	-1	0	1	0	0	1	3	3	
0	0	1	1	1	-1	0	1	0	0	1			
0	0	1	1	1	-1	0	1	0	0	1			
0	0	1	1	1									
0	0	1	1	1									

- Same K , new patch, new number. **The weights did not change**, only where we look

Worked example: slide once more

input X					kernel K			patch $\times$ kernel			output Y (sum = 0)		
0	0	1	1	1	-1	0	1	-1	0	1	3	3	0
0	0	1	1	1	-1	0	1	-1	0	1			
0	0	1	1	1	-1	0	1	-1	0	1			
0	0	1	1	1	-1	0	1	-1	0	1			
0	0	1	1	1	-1	0	1	-1	0	1			

- The patch is now all 1s, so each row gives $-1 + 0 + 1 = 0$: **no edge here**, and the output is 0
- That is the end of the first row. The window then drops down one row and repeats

Worked example: every position

input X					kernel K			output Y		
0	0	1	1	1	-1	0	1	3	3	0
0	0	1	1	1	-1	0	1	3	3	0
0	0	1	1	1	-1	0	1	3	3	0
0	0	1	1	1	-1	0	1	3	3	0
0	0	1	1	1	-1	0	1	3	3	0

- Big where the patch **looks like the kernel** (dark-to-bright, left to right), zero where the patch is flat
- The output is 3×3 : a 3×3 window only fits in 3 places across a row of 5

Output size

Input $n_h \times n_w$, **kernel** $k_h \times k_w$


$$(n_h - k_h + 1) \times (n_w - k_w + 1)$$

- The window has to fit entirely inside the input

Our example


- $n = 5, k = 3$
- $5 - 3 + 1 = 3$, so 3×3 out

A single filter on a real image



- **Red:** brighter on the right of the window. **Blue:** brighter on the left. **Gray:** flat
- The filter lights up on the **vertical strokes**, and only at their edges

What the filter sees



- Different kernel, different pattern detected, **same sliding procedure**
- A filter is a **pattern detector**: its output map says *where* in the image the pattern appears

Same shape, different name

A linear layer, one output

$$y = \sum_i w_i x_i + b$$

- Weights multiply **every** input
- One set of weights, used once

A convolution, one output position

$$Y[i, j] = \sum_{a, b} K[a, b] X[i+a, j+b] +$$

- Weights multiply **one patch**
- The *same* weights are reused at **every** position

Far fewer parameters

Map a 28 x 28 image to six 26 x 26 feature maps:

fully connected layer: 3,183,960 parameters

convolution, 3x3, 6 filters: 60 parameters

- **Locality:** each output looks at 9 pixels, not 784
- **Weight sharing:** the same 9 weights are reused at every position
- The kernel values are **not designed by hand**. They are parameters, found by gradient descent

Padding and stride

Controlling the output size

The output keeps shrinking

- Every convolution loses $k - 1$ rows and columns: $28 \rightarrow 26 \rightarrow 24 \rightarrow \dots$
- Pixels at the **edge** appear in fewer windows than pixels in the middle
- Sometimes we also want to shrink the output *on purpose*, to look at the image more coarsely

Two knobs: **padding** and **stride**

Padding: add a border of zeros

input with padding $p = 1$

0	0	0	0	0	0	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	0	0	0	0

output: 5×5

0	2	2	0	-2
0	3	3	0	-3
0	3	3	0	-3
0	3	3	0	-3
0	2	2	0	-2

- Padding p adds p rows and columns of zeros on **every side**
- With $k = 3$ and $p = 1$, the output is the same size as the input

Stride: skip positions

- **Stride** s is how far the window moves each step. So far, $s = 1$
- With $s = 2$, the window jumps two pixels: we compute roughly **half** as many outputs per direction
- Larger stride means a smaller, coarser output map

$s = 1$: windows start at columns 0, 1, 2

$s = 2$: windows start at columns 0, 2

Output size with padding and stride

$$\left\lfloor \frac{n + 2p - k}{s} \right\rfloor + 1$$

- n : input size, k : kernel size, p : padding per side, s : stride
- Apply it to the height and the width separately
- $\lfloor \cdot \rfloor$ rounds down: a window that would hang off the edge is simply not used

Worked example: padding 1, stride 2

padded input (7 × 7), windows jump 2 columns

0	0	0	0	0	0	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	1	1	1	0
0	0	0	0	0	0	0

output (3 × 3)

0	2	-2
0	3	-3
0	2	-2

- Padding adds the gray border; stride 2 moves the window two columns (and two rows) at a time
- The first row of the output comes from the three windows shown. Windows overlap by one column

Worked example: our 5×5 input, 3×3 kernel

Padding p	Stride s	$\lfloor (5 + 2p - 3)/s \rfloor + 1$	Output
0	1	$\lfloor 2/1 \rfloor + 1$	3×3
1	1	$\lfloor 4/1 \rfloor + 1$	5×5
0	2	$\lfloor 2/2 \rfloor + 1$	2×2
1	2	$\lfloor 4/2 \rfloor + 1$	3×3

Coming up: this is one line in PyTorch

```
1 x = torch.zeros(1, 1, 5, 5)          # (batch, channels, height, width)
2
3 for p, s in [(0, 1), (1, 1), (0, 2), (1, 2)]:
4     conv = nn.Conv2d(1, 1, kernel_size=3, padding=p, stride=s)
5     print(f"padding={p}, stride={s} -> {tuple(conv(x).shape[2:])}")
```

```
padding=0, stride=1 -> (3, 3)
padding=1, stride=1 -> (5, 5)
padding=0, stride=2 -> (2, 2)
padding=1, stride=2 -> (3, 3)
```

- Same four sizes as the table: 3×3 , 5×5 , 2×2 , 3×3
- We will build up to `nn.Conv2d` shortly. For now, notice that `padding` and `stride` are just arguments


Channels

More than one filter, more than one input map

One filter, one pattern

- A single filter detects a single pattern
- Images have many kinds of patterns: vertical strokes, horizontal strokes, diagonals, curves
- So a layer uses **many filters at once**, each with its own weights, each producing its own output map

Multiple filters: multiple output channels



- Each filter produces one **feature map**. Stacking them gives an output with several **channels**
 - In this case we output 4 channels
- These four are hand-made; in a real CNN the filter values are **learned**

Multiple input channels

- A color image already has 3 channels (red, green, blue). Later layers take the previous layer's feature maps as *their* input channels
- A kernel for a c_{in} -channel input has its own $k \times k$ slice **per input channel**
- Each slice is slid over its channel; the results are **added together** into one output value

Multiple input channels: the data flow

input, channel 1

0	0	1	1	1
0	0	1	1	1
0	0	1	1	1
0	0	1	1	1
0	0	1	1	1

channel 2

0	0	0	0	0
0	0	0	0	0
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

channel 3

0	0	1	0	0
0	1	0	0	0
1	0	0	0	0
0	0	0	0	0
0	0	0	0	0

kernel slices

-1	0	1
-1	0	1
-1	0	1

-1	-1	-1
0	0	0
1	1	1


1	-1	-1
-1	1	-1
-1	-1	1

partial sums

→ 3

→ 3

→ -1



ONE output value

Shapes in one convolutional layer

	Shape	What it is
Input	(c_{in}, H, W)	c_{in} channels, each an $H \times W$ grid
Weights	(c_{out}, c_{in}, k, k)	c_{out} filters, each with a $k \times k$ slice per input channel
Bias	(c_{out})	one number per filter
Output	(c_{out}, H', W')	one feature map per filter

- This describes **one** 2D convolutional layer. A CNN stacks several, and each layer's output becomes the next layer's input
- H', W' come from the output-size formula (kernel size, padding, stride)

Counting parameters

Start with one filter

- It has a $k \times k$ slice for each of the c_{in} input channels: $c_{in} \cdot k \cdot k$ weights
- It has a single bias: $+1$

A layer has c_{out} filters, each with its own weights and bias

$$c_{out} \cdot (c_{in} \cdot k \cdot k + 1) = \underbrace{c_{out} \cdot c_{in} \cdot k \cdot k}_{\text{weights}} + \underbrace{c_{out}}_{\text{biases}}$$

- c_{out} appears twice because every filter has its **own** bias

Worked examples

- 6 filters, 1 input channel, $k = 3$: $6 \cdot 1 \cdot 9 + 6 = 60$
- 16 filters, 6 input channels, $k = 3$: $16 \cdot 6 \cdot 9 + 16 = 880$

```
1 for c_in, c_out in [(1, 6), (6, 16)]:
2     conv = nn.Conv2d(c_in, c_out, kernel_size=3)
3     print(f"{c_in} -> {c_out} channels:", sum(p.numel() for p in conv.parameters()), "parameters")
```

1 -> 6 channels: 60 parameters

6 -> 16 channels: 880 parameters

- The count does **not** depend on the image size $H \times W$

Pooling

Summarizing a neighborhood

Exact position is too much detail

- After a filter, we know *where* it fired, down to the pixel
- But a “4” drawn one pixel to the left is still a 4
- These outputs (feature maps) are *big* and stacking more layers on top of them is expensive

The Big Idea

Keep the strongest response in each small neighborhood, and forget exactly where inside it that response was.

Max pooling example

feature map (4×4)

1	3	2	0
4	2	1	5
0	1	6	2
3	2	1	1


2×2 max pool

4	5
3	6


- Window 2×2 , stride 2 (non-overlapping): take the **max** of each block

Pooling on a real feature map

feature map (after ReLU): 26×26



after 2×2 max pool: 13×13



- Same pattern, a quarter of the numbers

What pooling keeps and discards

Keeps

- *Whether* the pattern appeared in the neighborhood
- How strong the strongest response was

Discards

- *where* in the neighborhood it was
- Everything weaker than the maximum

-
- Pooling has **no parameters**; there is nothing to learn
 - Output size follows the same formula as convolution (use the pool size as k)

Building convolution layers

From the formula to `nn.Module`

Step 1: the operation from scratch

```
1 def corr2d(X, K):
2     h, w = K.shape
3     Y = torch.zeros(X.shape[0] - h + 1, X.shape[1] - w + 1)
4     for i in range(Y.shape[0]):
5         for j in range(Y.shape[1]):
6             Y[i, j] = (X[i:i + h, j:j + w] * K).sum()
7     return Y
8
9 X = torch.tensor([[0., 0, 1, 1, 1]] * 5)
10 K = torch.tensor([[-1., 0, 1]] * 3)
11 corr2d(X, K)
```

```
tensor([[3., 3., 0.],
        [3., 3., 0.],
        [3., 3., 0.]])
```

- Same numbers as our hand-worked example: every row [3, 3, 0]

Step 2: a layer is a Module

```
1 class Conv2D(nn.Module):
2     def __init__(self, kernel_size):
3         super().__init__()
4         self.weight = nn.Parameter(torch.rand(kernel_size))
5         self.bias = nn.Parameter(torch.zeros(1))
6
7     def forward(self, X):
8         return corr2d(X, self.weight) + self.bias
```

- Same interface as every model so far: `__init__` creates parameters, `forward` is the recipe
- The kernel and bias are `nn.Parameters`, so the optimizer will find them

Step 3: the kernel can be learned

Hide a kernel in some data, then see if training finds it.

```
1 x = torch.tensor(digit)
2 target = corr2d(x, K)          # "what the right filter would produce"
3
4 layer = Conv2D((3, 3))
5 opt = torch.optim.SGD(layer.parameters(), lr=0.5)
6 for step in range(300):
7     loss = ((layer(x) - target) ** 2).mean()
8     opt.zero_grad()
9     loss.backward()
10    opt.step()
11
12 print(layer.weight.detach().round(decimals=1))
```

tensor([[-0.9000, -0.1000, 1.1000],
 [-1.0000, 0.1000, 0.9000],
 [-1.0000, 0.0000, 1.0000]])

- Starting from random values, the training loop from last lecture recovers K : left column ≈ -1 , middle ≈ 0 , right ≈ 1
- Nobody designs the filters in a CNN. We choose how many, and how big; **training chooses the values**

The built-in version: `nn.Conv2d`

PyTorch ships the same operation. Does it agree with ours?

```
1 ours = Conv2D((3, 3))
2 theirs = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3)
3
4 with torch.no_grad():                # give both the same weights
5     theirs.weight.copy_(ours.weight.reshape(1, 1, 3, 3))
6     theirs.bias.copy_(ours.bias)
7
8 x = torch.tensor(digit)
9 print(torch.allclose(ours(x), theirs(x[None, None])[0, 0]))
```

True

- Same operation, same numbers. `nn.Conv2d` replaces our `Conv2D`; it is not a new idea
- What it adds: many input and output channels, batches, `padding` and `stride`, and speed (no Python loops)

nn.Conv2d: shapes and arguments

```
1 conv = nn.Conv2d(in_channels=1, out_channels=6, kernel_size=3)
2
3 print("weight:", tuple(conv.weight.shape))    # (c_out, c_in, k, k)
4 print("bias:  ", tuple(conv.bias.shape))      # (c_out,)
5
6 batch = torch.zeros(8, 1, 28, 28)           # (batch, channels, height, width)
7 print("output:", tuple(conv(batch).shape))
```

```
weight: (6, 1, 3, 3)
bias:   (6,)
output: (8, 6, 26, 26)
```

- PyTorch layers expect **batches** of shape (batch, channels, height, width)
- Padding and stride are arguments: `nn.Conv2d(1, 6, 3, padding=1, stride=2)`
- `nn.MaxPool2d(2)` is the pooling layer: pool size 2, stride 2 by default

Stacking into a small CNN

```
1 net = nn.Sequential(  
2     nn.Conv2d(1, 6, kernel_size=3), nn.ReLU(), nn.MaxPool2d(2),  
3     nn.Conv2d(6, 16, kernel_size=3), nn.ReLU(), nn.MaxPool2d(2),  
4     nn.Flatten(),  
5     nn.Linear(16 * 5 * 5, 10),  
6 )  
7  
8 X = torch.zeros(1, 1, 28, 28)  
9 for layer in net:  
10     X = layer(X)  
11     print(f"{type(layer).__name__:10s} -> {tuple(X.shape)}")
```

```
Conv2d      -> (1, 6, 26, 26)  
ReLU        -> (1, 6, 26, 26)  
MaxPool2d   -> (1, 6, 13, 13)  
Conv2d      -> (1, 16, 11, 11)  
ReLU        -> (1, 16, 11, 11)  
MaxPool2d   -> (1, 16, 5, 5)  
Flatten     -> (1, 400)  
Linear      -> (1, 10)
```

Tracing shapes: why the 400?

- $28 \rightarrow 26$ (conv, no padding) $\rightarrow 13$ (pool) $\rightarrow 11$ (conv) $\rightarrow 5$ (pool, rounds down)
- The last feature maps are $16 \times 5 \times 5$, so **Flatten** gives $16 \cdot 5 \cdot 5 = 400$ numbers
- That is where **`nn.Linear(400, 10)`** comes from.
 - **Shape mismatches here are the most common CNN bug: trace the shapes by hand first**

```
1 print(f"total parameters: {sum(p.numel() for p in net.parameters()):,}")
```

```
total parameters: 4,950
```

Throwback: the same input through an MLP

```
1 mlp = nn.Sequential(nn.Flatten(), nn.Linear(784, 16), nn.ReLU(), nn.Linear(16, 10))
2
3 def count(model):
4     return sum(p.numel() for p in model.parameters())
5
6 print(f"MLP, one tiny hidden layer (16 units): {count(mlp):,} parameters")
7 print(f"small CNN from the last slide: {count(net):,} parameters")
```

MLP, one tiny hidden layer (16 units): 12,730 parameters
small CNN from the last slide: 4,950 parameters

- Even a *tiny* MLP has more parameters than our CNN: its first layer alone is $784 \cdot 16 + 16 = 12,560$
- The CNN gets away with fewer because each 3×3 filter is **reused at every position**, instead of every pixel getting its own weight to every unit

Reference: a CNN is just another `nn.Module`


```
1 class SmallCNN(nn.Module):
2     def __init__(self, num_classes=10):
3         super().__init__()
4         self.network = nn.Sequential(
5             nn.Conv2d(1, 6, kernel_size=3), nn.ReLU(), nn.MaxPool2d(2),
6             nn.Conv2d(6, 16, kernel_size=3), nn.ReLU(), nn.MaxPool2d(2),
7             nn.Flatten(),
8             nn.Linear(16 * 5 * 5, num_classes),
9         )
10
11     def forward(self, X):
12         return self.network(X)
13
14 model = SmallCNN()
15 model(torch.zeros(8, 1, 28, 28)).shape
torch.Size([8, 10])
```

- Same `__init__` and `forward` pattern as the MLP. Training works exactly the same way: next lecture

Looking ahead: representations

What are the convolutional layers *for*?

The model has two parts



- The last layer is the **same linear classifier** we already know
- It does not see pixels. It sees the 400 numbers the convolutional layers produced

Foreshadowing: the convolutional layers produce features

- Each of those 400 numbers answers a question about the image: *how strongly is this pattern here?*
- The pattern detectors are **learned**, not hand-designed. Training picks the filters that make the classifier's job easiest
- Stack layers, and later filters look at patterns **of patterns**: a 3×3 filter on top of a 3×3 filter effectively sees a 5×5 region

The Big Idea (for the rest of the course)

A deep model learns to turn raw data into a representation, a set of features, that makes the final task easy.

- CNNs are our first example, and we will keep coming back to it: **what does the model learn to see?**

Summary

Recap: the convolution toolkit

Piece	What it does	Output size / cost
Convolution	Slide a shared kernel over the image; sum products at each position	$\lfloor (n + 2p - k) / s \rfloor + 1$
Padding p	Add a border of zeros so edges count and size is kept	grows the input by $2p$
Stride s	Move the window more than one pixel per step	shrinks output by $\approx s$
Channels	Many filters per layer; each input channel gets its own kernel slice	$c_{out} \cdot c_{in} \cdot k^2 + c_{out}$ parameters

Recap: key ideas

- **Locality and weight sharing:** look at small patches with the same weights everywhere
- A filter is a **pattern detector**; its output map says where the pattern appears
- Filter values are **learned** by gradient descent, just like every other weight
- A CNN layer is an `nn.Module` like any other: `nn.Conv2d`, `nn.MaxPool2d`, and `nn.Sequential`
- **Next time:** train a CNN end to end and compare it to an MLP on the same data