

Linear Regression for Classification

Advanced Machine Learning

Sam Ogden

DRAFT

This slide deck is still under active development. Expect changes before it is finalized.

Learning objectives

- Give categories their own model output, via one-hot labels and softmax
- Define cross-entropy loss for that output
- Show its gradient looks just like linear regression's

Classification

The Big Question

Predicting *which*, not *how much*

What we have vs. what's different

What we already have

- A dataset, a model, a loss
- A way to train it — all covered previously
- All of it assumed the answer is a single number

What's different now

- The answer is a **category**
 - spam or not spam
 - cat or dog
 - which digit
- No natural “how far off” between categories, the way there is between numbers

The Big Idea

Some questions don't have a numeric answer — they have a category as the answer.

A concrete example we'll carry through

We'll sort pieces of fruit into **apple**, **banana**, or **orange**, from two features:

weight (g)	color score (0=green, 1=red)	fruit
150	0.9	apple
120	0.2	banana
180	0.7	orange

Why not just reuse linear regression directly?

Categories aren't numbers

- We're used to fitting a line to numbers
- Apple, banana, orange aren't numbers — there's no natural “less than” or “far apart” between them
- Coding them as 0, 1, 2 invents structure that isn't really there

The Big Idea

Each *individual* category is binary — it is, or it isn't.

Let's just try it anyway

- Code the labels as integers: apple = 0, banana = 1, orange = 2
- Fit the usual linear model against those codes:

$$\hat{y} = \mathbf{w} \cdot \mathbf{x} + b$$

- Predict a class by rounding $\hat{y}$ to the nearest code

Two problems with this

- **We invented an ordering** — this encoding says banana (1) sits *between* apple (0) and orange (2); relabel the fruits differently and the model's behavior changes
- **The output doesn't match the question** — $\hat{y}$ can be anything (-1.3, 4.7, ...); rounding to the nearest code is a patch, not a fix, and a 1.99 vs. a 1.01 both round to “banana” without feeling equally confident

What we actually need

- A way to write “which category” without implying order or distance → **one-hot encoding**
- A model output that’s bounded and behaves like a confidence → **softmax**

One-hot encoding

The Big Idea

Give every category its own slot, and just mark which one is true.

Our fruit, one-hot

fruit	apple	banana	orange
apple	1	0	0
banana	0	1	0
orange	0	0	1

Note

Compare to the integer coding from a few slides ago — apple, banana, and orange no longer have any numeric relationship to each other at all

Written out

- K categories $\rightarrow$ a vector of length K , one slot per category
- Slot j is 1 if the example belongs to category j , 0 otherwise:

$$y_j = \begin{cases} 1 & \text{example is class } j \\ 0 & \text{otherwise} \end{cases}$$

- Exactly one slot is 1 — that’s the “hot” one — and its *position* carries all the information; nothing else in the vector does

A “linear” model, per class

Linear scores — for now

One score per class (recap)

- One shared model, not K separate ones — $\mathbf{W}$ has one row per class, $\mathbf{b}$ one entry per class:

$$\mathbf{o} = \mathbf{W}\mathbf{x} + \mathbf{b}, \quad o_j = \mathbf{w}_j \cdot \mathbf{x} + b_j$$

- $\mathbf{o}$ is just this layer's raw output (hence the letter) — same forward pass as multi-output regression, covered previously
- Each o_j is called a **logit**: the model's raw, unnormalized opinion about class j — it can be negative, has no upper bound, and the o_j 's don't sum to anything meaningful

⚠ Important

$\mathbf{o}$ is not our prediction yet — one more function still needs to act on it

Worked example: our fruit

Suppose, for one particular (weight, color) example, evaluating $o_j = \mathbf{w}_j \cdot \mathbf{x} + b_j$ for each class gives:


apple	banana	orange
-0.3	2.1	0.6

- Banana has the highest logit, so *right now* the model favors banana

Warning

So the answer is just “banana,” done? That’s not actually what we do — why not?

The picture so far



- $x \rightarrow$ linear layer $\rightarrow o$ is the whole story so far
- Something still needs to turn o into an actual prediction — that's next

Softmax

From scores to probabilities

The Big Idea


We make each prediction *act* like a probability – without changing which class wins.

The softmax function

$$\hat{y}_j = \text{softmax}(\mathbf{o})_j = \frac{e^{o_j}}{\sum_{k=1}^K e^{o_k}}$$

- Takes in all K logits at once, returns one probability per class
- $\hat{\mathbf{y}}$ is now our model's actual prediction — a full probability distribution over the classes

The full picture



- Two functions, applied one after the other: a **linear** layer, then **softmax**
- o was never the answer — it's an intermediate value that softmax turns into $\hat{y}$

Breakdown: numerator and denominator

Numerator

$$e^{o_j}$$

- Exponentiate *this* class's logit
- Always positive
- Same ranking as o_j

Denominator

$$\sum_{k=1}^K e^{o_k}$$

- The *same kind of term*
- Summed over every class
- Includes j itself

Note

Divide one by the other: squeezes every class's share into $[0, 1]$, and all K shares add up to exactly 1

So what does this actually do?

- It's not just squashing into $[0, 1]$ — exponentiating also **magnifies** the gaps between logits
- Our banana/orange logits differed by only 1.5; their probabilities differ by a factor of $e^{1.5} \approx 4.5$
- A modest lead in logits turns into a much larger lead in confidence

Worked example: our fruit

Recall our logits: apple = -0.3 , banana = 2.1 , orange = 0.6

	apple	banana	orange
e^{o_j}	0.74	8.17	1.82
$\hat{y}_j$	0.07	0.76	0.17

- Same winner as before (banana) — softmax never changes who wins
- But now we also know *how confident*: about 76%, not just “highest”

The special case: two classes

- With only $K = 2$ classes, softmax simplifies to:

$$\hat{y}_1 = \frac{1}{1 + e^{-(o_1 - o_2)}}$$

- That's the **sigmoid** function, applied to the *difference* between the two logits
- This special case has its own name — **logistic regression** — softmax is its generalization to more than two classes

What softmax buys us

- **A real probability distribution:** non-negative, sums to 1, ranks classes the same way the logits did
- **A smooth function:** softmax is differentiable, unlike `argmax`
 - That matters a lot! To train with gradient descent, we need a loss we can actually take a gradient of

Loss: Cross entropy

Loss function for multi-classification

Why not just reuse MSE?

- MSE only knows “off by how much”; it can’t tell *unsure* apart from *confidently wrong*
- Being *confidently* wrong should cost a lot more than being unsure

The Big Idea

Penalize the model by how little probability it gave the correct answer.

Cross-entropy loss

$$L(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{j=1}^K y_j \log \hat{y}_j$$

- $\mathbf{y}$: the one-hot true label
- $\hat{\mathbf{y}}$: the softmax output — our predicted probabilities
- One number, same as MSE was — small if we're right, big if we're wrong

Only one term survives

- y_j is 0 for every class except the true one, call it t
- Every term in the sum vanishes except $j = t$
- The whole sum collapses to just:

$$L = -\log \hat{y}_t$$

- Cross-entropy only ever looks at the probability the model gave the *correct* class


Breakdown: why the negative log?

$\hat{y}_t$ (true class probability)	$-\log \hat{y}_t$
0.99	0.01
0.76	0.27
0.50	0.69
0.10	2.30
0.01	4.61

- Confident *and* correct ($\hat{y}_t$ near 1) $\rightarrow$ loss near 0
- Confident *and* wrong ($\hat{y}_t$ near 0) $\rightarrow$ loss blows up
- Being unsure costs a little; being confidently wrong costs a lot

Log-loss, visualized

One classification act — one example, one true class — costs $-\log \hat{y}_t$:



- The curve is steep near 0, flat near 1 — the same asymmetry as the table, just visible now
- The five orange points are exactly the five rows from the table

Worked example: our fruit

Recall our softmax output: apple = 0.07, banana = 0.76, orange = 0.17 — and the true label is banana

- $L = -\log(0.76) \approx 0.27$
- Had the model instead favored apple, with only 0.07 on banana:
 $L = -\log(0.07) \approx 2.66$
- Same “wrong or right” call either way — cross-entropy is what tells them apart

Differentiating the loss

Surprisingly easy calculus

The Big Idea

This scary-looking loss has a gradient just as simple as linear regression's.

Rewrite the loss and softmax first

Recall $L = -\log(\hat{y}_t)$, where $\hat{y}_t = \frac{e^{o_t}}{\sum_k e^{o_k}}$ — substituting one into the other **fuses log-loss and softmax into a single expression**, before we differentiate anything:

$$-\log \frac{e^{o_t}}{\sum_{k=1}^K e^{o_k}} = -\log(e^{o_t}) + \log \sum_{k=1}^K e^{o_k} = -o_t + \log \sum_{k=1}^K e^{o_k}$$

- Step 1: $\log(a/b) = \log a - \log b$
- Step 2: $\log(e^{o_t}) = o_t$ — true because log here means the **natural log** (ln), not $\log_{10}$
- Still just algebra — no derivative yet. We're rewriting L into a form that's easy to differentiate next

Take the derivative

- Differentiate term by term with respect to o_j :

$$\frac{\partial L}{\partial o_j} = -y_j + \frac{e^{o_j}}{\sum_{k=1}^K e^{o_k}}$$

- The $-o_t$ term only depends on o_j when $j = t$ — its derivative is -1 then, 0 otherwise, which is exactly $-y_j$
- The second term is just softmax again, $\hat{y}_j$ leaving us with:

$$\frac{\partial L}{\partial o_j} = \hat{y}_j - y_j$$

Breakdown: the same shape as before

- Regression's gradient, covered previously: $\frac{\partial L}{\partial w} = (\hat{y}_i - y_i) x_i$ — “prediction minus truth, times input”
- Same shape here, just once per class:

$$\frac{\partial L}{\partial \mathbf{w}_j} = (\hat{y}_j - y_j) \mathbf{x}, \quad \frac{\partial L}{\partial b_j} = \hat{y}_j - y_j$$

! Important

All that softmax and log-loss machinery, and the gradient still just says “how wrong were we” — negative for the true class (push its logit up), positive for every other class (push them down)

Putting it together: training a classifier

Same loop, new loss

The Big Idea

**Nothing about the training loop changes
— only the loss, and what “prediction”
means.**

What changed / what didn't

Same as before

- The training loop itself — sample, gradient, step, repeat
- The learning rate η
- $(\mathbf{W}, \mathbf{b})$ as the trainable parameters

New this lecture

- Prediction: a probability vector, via softmax
- Loss: cross-entropy, not MSE
- Gradient: $\hat{y}_j - y_j$ — same shape as before, just $\hat{y}$ is now *computed* differently (via softmax, not a plain dot product)

Summary

The recipe

Task	classify each example into one of K categories
Example	apple / banana / orange, from weight + color
Model output	logits $\rightarrow$ softmax $\rightarrow$ a probability vector
Loss	cross-entropy

Note

This recipe — categorical output, softmax, cross-entropy — is the template for classification problems in general, not just this one

Tip

In lab: walk one example through a full forward pass (logits $\rightarrow$ softmax $\rightarrow$ cross-entropy) and backward pass (the gradient) **by hand**, with actual numbers